

# Milestone3-Task3

April 16, 2022

## 1 Task 3

## 2 Imports

```
[1]: import numpy as np
import pandas as pd
from joblib import dump, load
from sklearn.metrics import mean_squared_error
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split, cross_validate
import matplotlib.pyplot as plt
plt.style.use('ggplot')
plt.rcParams.update({'font.size': 16, 'axes.labelweight': 'bold', 'figure.
↳figsize': (8,6)})
## add any other additional packages that you need. You are free to use any
↳packages for vizualization.
import altair as alt
alt.data_transformers.disable_max_rows()
```

```
[1]: DataTransformerRegistry.enable('default')
```

### 2.1 Part 1:

Recall as a final goal of this project. We want to build and deploy ensemble machine learning models in the cloud, where features are outputs of different climate models and the target is the actual rainfall observation. In this milestone, you'll actually build these ensemble machine learning models in the cloud.

#### Your tasks:

1. Read the data CSV from your s3 bucket.
2. Drop rows with nans.
3. Split the data into train (80%) and test (20%) portions with `random_state=123`.
4. Carry out EDA of your choice on the train split.
5. Train ensemble machine learning model using `RandomForestRegressor` and evaluate with metric of your choice (e.g., RMSE) by considering `Observed` as the target column.
6. Discuss your results. Are you getting better results with ensemble models compared to the individual climate models?

Recall that individual columns in the data are predictions of different climate models.

```
[2]: ## Depending on the permissions that you provided to your bucket you might need
      ↳ to provide your aws credentials
      ## to read from the bucket, if so provide with your credentials and pass as
      ↳ storage_options=aws_credentials
      aws_credentials = {"key" : "ASIARS6J4JEAH2ZOBKCD",
                        "secret": "vhJG5S0k783tWBCqOZlkY8NTwaXZl1jMtfq8gGjT",
                        "token": "FwoGZXIvYXZELj//////////
      ↳ wEaDLBw9pd5Kse5Pi7LiyLCAasr9ifix3bKDd6nf5hHHCxKTYtX7shxFfRM73uJApDrJmPlhXN1bQ1swlqSJjGrGoly
      ↳ waunsP2WlwU2ygC3iiBz2/
      ↳ y1d3UhV+67jqnXdkUQIxVy8nNOD0v7u6IErx8VlqzbTynpMODyXeuBTJzLlVWhQtPQ8chuQIIR9aFi+hpKg28x+0mbV
                        }

      #df = pd.read_csv("s3://mds-s3-group18/output/ml_data_SYD.csv",
      ↳ storage_options=aws_credentials, index_col=0, parse_dates=True)
      df = pd.read_csv("ml_data_SYD.csv")
```

```
[3]: df
```

```
[3]:
```

	time	ACCESS-CM2_rainfall	ACCESS-ESM1-5_rainfall	\
0	1889-01-01	0.040427	1.814552	
1	1889-01-02	0.073777	0.303965	
2	1889-01-03	0.232656	0.019976	
3	1889-01-04	0.911319	13.623777	
4	1889-01-05	0.698013	0.021048	
...	...	...	...	
46015	2014-12-27	0.033748	0.123476	
46016	2014-12-28	0.094198	2.645496	
46017	2014-12-29	0.005964	3.041667	
46018	2014-12-30	0.000028	1.131412	
46019	2014-12-31	0.532747	2.370896	
	AWI-ESM-1-1-LR_rainfall	BCC-CSM2-MR_rainfall	BCC-ESM1_rainfall	\
0	3.557934e+01	4.268112e+00	1.107466e-03	
1	4.596520e+00	1.190141e+00	1.015323e-04	
2	5.927467e+00	1.003845e-09	1.760345e-05	
3	8.029624e+00	8.225225e-02	1.808932e-01	
4	2.132686e+00	2.496841e+00	4.708019e-09	
...	...	...	...	
46015	1.451179e+00	3.852845e+01	2.061717e-03	
46016	4.249335e+01	5.833801e-01	5.939502e-09	
46017	2.898325e+00	9.359547e-02	2.000051e-08	
46018	2.516381e-01	1.715028e-01	7.191735e-05	
46019	1.047835e-13	4.437736e+00	2.863683e-01	
	CMCC-CM2-HR4_rainfall	CMCC-CM2-SR5_rainfall	CMCC-ESM2_rainfall	\

0	1.141054e+01	3.322009e-08	2.668800
1	4.014984e+00	1.312700e+00	0.946211
2	9.660565e+00	9.103720e+00	0.431999
3	3.951528e+00	1.317160e+01	0.368693
4	2.766362e+00	1.822940e+01	0.339267
...	...	...	...
46015	8.179260e-09	1.171263e-02	0.090786
46016	8.146937e-01	4.938899e-01	0.000000
46017	2.532205e-01	1.306046e+00	0.000002
46018	8.169252e-02	1.722262e-01	0.788577
46019	6.343592e+00	6.368303e-01	0.442130

	CanESM5_rainfall	...	MPI-ESM-1-2-HAM_rainfall	\
0	1.321215	...	4.244226e-13	
1	2.788724	...	4.409552e+00	
2	0.003672	...	2.269300e-01	
3	0.013578	...	2.344586e-02	
4	0.002468	...	4.270161e-13	
...	...	...	...	
46015	59.895053	...	4.726998e-13	
46016	0.512632	...	4.609420e-13	
46017	37.169669	...	2.016156e+01	
46018	7.361246	...	9.420543e+00	
46019	0.306608	...	1.031899e+01	

	MPI-ESM1-2-HR_rainfall	MPI-ESM1-2-LR_rainfall	MRI-ESM2-0_rainfall	\
0	1.390174e-13	6.537884e-05	3.445495e-06	
1	1.222283e-01	1.049131e-13	4.791993e-09	
2	3.762301e-01	9.758706e-14	6.912302e-01	
3	4.214019e-01	7.060915e-03	3.835721e-02	
4	1.879692e-01	4.504985e+00	3.506923e-07	
...	...	...	...	
46015	1.326889e-01	1.827857e+00	6.912632e-03	
46016	1.644482e+00	7.242920e-01	2.836752e-03	
46017	1.506439e+00	1.049481e-01	8.137182e+00	
46018	6.242895e+00	1.245115e-01	9.305263e-03	
46019	4.765813e+01	3.274323e-01	6.854985e-11	

	NESM3_rainfall	NorESM2-LM_rainfall	NorESM2-MM_rainfall	\
0	1.576096e+01	4.759651e-05	2.451075	
1	3.675510e-01	4.350863e-01	0.477231	
2	1.562869e-01	9.561101e+00	0.023083	
3	2.472226e-07	5.301038e-01	0.002699	
4	1.949792e-13	1.460928e-10	0.001026	
...	...	...	...	
46015	2.171327e-03	1.620489e+00	2.084252	
46016	1.344768e+01	2.391159e+00	1.644527	

46017	2.547820e+01	1.987695e-12	0.205036
46018	4.192948e+00	2.150346e+00	0.000017
46019	2.067847e+00	2.349716e+01	0.035319

	SAMO-UNICON_rainfall	TaiESM1_rainfall	observed_rainfall
0	0.221324	2.257933	0.006612
1	3.757179	2.287381	0.090422
2	0.253357	1.199909	1.401452
3	2.185454	2.106737	14.869798
4	2.766507	1.763335	0.467628
...	...	...	...
46015	0.868046	17.444923	0.037472
46016	0.782258	1.569647	0.158061
46017	2.140723	1.444630	0.025719
46018	29.714692	0.716019	0.729390
46019	59.724062	3.240185	0.008076

[46020 rows x 27 columns]

```
[4]: ## Use your ML skills to get from step 1 to step 6
```

```
[5]: # Drop rows with na
df = df.dropna()
```

```
[6]: # Split the data into train (80%) and test (20%) portions with random_state=123

train, test = train_test_split(df, test_size=0.2, random_state=123)
```

```
[7]: # Carry out EDA of your choice on the train split

# Distribution of target label
(alt.Chart(train).mark_bar().encode(
    alt.X('observed_rainfall', bin=alt.Bin(maxbins=10), title='Observed_
↳ Rainfall (mm)'),
    y='count()'
))
```

```
[7]: alt.Chart(...)
```

```
[8]: # Observed rainfall over time
train['time'] = pd.to_datetime(train['time'])
(alt.Chart(train).mark_line().encode(
    x='time',
    y='observed_rainfall'
))
```

```
[8]: alt.Chart(...)
```

```
[9]: # Train ensemble machine learning model using RandomForestRegressor and
      ↪ evaluate with metric of your choice (e.g., RMSE) by considering Observed as
      ↪ the target column
y_train = train['observed_rainfall']
X_train = train.drop(['observed_rainfall', 'time'], axis=1)

m_rf = RandomForestRegressor()
res = cross_validate(m_rf, X_train, y_train, cv=5,
                    ↪ scoring='neg_root_mean_squared_error')
print(f"CV-train RMSE: {res['test_score'].mean()*-1}")
m_rf.fit(X_train, y_train)
```

CV-train RMSE: 8.330195567524267

```
[9]: RandomForestRegressor()
```

```
[10]: m_rf.fit(X_train, y_train)
```

```
[10]: RandomForestRegressor()
```

```
[11]: # Predict on test dataset
y_test = test['observed_rainfall']
X_test = test.drop(['observed_rainfall', 'time'], axis=1)

test['Ensemble'] = m_rf.predict(X_test)
```

```
[12]: # Calculate RMSE for all models
test_wide = test.drop(['time', 'observed_rainfall'], axis=1)
test_wide = pd.melt(test_wide, var_name='model', value_name='rainfall')

test_result = {}

for model in test_wide['model'].unique():
    pred = test_wide.query('model == @model')['rainfall']
    test_result[model] = mean_squared_error(y_test, pred, squared=False)
```

```
[13]: pd.DataFrame([test_result]).T.rename(columns={0: "RMSE"}).sort_values('RMSE')
```

```
[13]:
```

	RMSE
Ensemble	8.846681
KIOST-ESM_rainfall	9.600480
FGOALS-g3_rainfall	9.687788
MRI-ESM2-0_rainfall	9.922795
MPI-ESM1-2-HR_rainfall	9.969823
NESM3_rainfall	9.978137
MPI-ESM1-2-LR_rainfall	10.260886
NorESM2-LM_rainfall	10.410145

EC-Earth3-Veg-LR_rainfall	10.453606
GFDL-CM4_rainfall	10.511682
BCC-ESM1_rainfall	10.615578
CMCC-CM2-HR4_rainfall	10.643204
ACCESS-ESM1-5_rainfall	10.695305
BCC-CSM2-MR_rainfall	10.761381
MPI-ESM-1-2-HAM_rainfall	10.932004
NorESM2-MM_rainfall	10.939740
AWI-ESM-1-1-LR_rainfall	10.996616
ACCESS-CM2_rainfall	11.038999
CanESM5_rainfall	11.151318
CMCC-ESM2_rainfall	11.246493
MIROC6_rainfall	11.352976
INM-CM4-8_rainfall	11.451635
CMCC-CM2-SR5_rainfall	11.480614
TaiESM1_rainfall	11.528083
SAMO-UNICON_rainfall	11.678749
INM-CM5-0_rainfall	12.250223

Discuss your results. Are you getting better results with ensemble models compared to the individual climate models

Yes, we are getting better results with ensemble model compared to individual climate model. This is because by ensembling different models together, we are able to cancel out the errors made by different models, and combine different mapping function learnt by each individual climate model.

## 2.2 Part 2:

### 2.2.1 Preparation for deploying model next week

*NOTE: Complete task 4 from the milestone3 before coming here*

We've found the best hyperparameter settings with MLlib (from the task 4 from milestone3), here we then use the same hyperparameters to train a scikit-learn model.

```
[14]: model = RandomForestRegressor(n_estimators=100, max_depth=5, bootstrap=False)
      model.fit(X_train, y_train)
```

```
[14]: RandomForestRegressor(bootstrap=False, max_depth=5)
```

```
[15]: print(f"Train RMSE: {mean_squared_error(y_train, model.predict(X_train),
      ↪squared=False):.2f}")
      print(f" Test RMSE: {mean_squared_error(y_test, model.predict(X_test),
      ↪squared=False):.2f}")
```

Train RMSE: 7.91

Test RMSE: 8.71

```
[16]: # ready to deploy  
      dump(model, "model.joblib")
```

```
[16]: ['model.joblib']
```

*Upload model.joblib to s3 under output folder. You choose how you want to upload it (using CLI, SDK, or web console).*